



Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Processamento e Recuperação de Informação

Efficient Similarity Search

Departamento de Engenharia Informática
Instituto Superior Técnico

1º Semestre
2018/2019



Outline

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- 1 Finding Similar Items
- 2 Shingles
- 3 Minhashing
- 4 Locality-sensitive hashing



Bibliography

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Jure Leskovec, Anand Rajaraman, and Jeff Ullman, Mining of Massive Datasets, Chapter 3



High Dimensional Data

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Many real-world problems

- Web Search and Text Mining
 - Billions of documents, millions of terms
- Product Recommendations
 - Millions of customers, millions of products
- Scene Completion, other graphics problems
 - Image features
- Online Advertising, Behavioral Analysis
 - Customer actions (e.g., websites visited, searches)



A Common Metaphor

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Many problems can be expressed as finding **similar** sets.

Find near-neighbors in high-dimensional space.

Examples:

- Pages with similar words
 - For duplicate detection, classification by topic
- Customers who purchased similar products
 - NetFlix users with similar tastes in movies
- Products with similar customer sets
- Images with similar features
- Users who visited the similar websites



Distance Measures

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

We formally define **near neighbors** as points that are a **small distance** apart.

For each use case, we need to define what distance means.

Two major classes of distance measures:

- A Euclidean distance is based on the locations of points in such a space
- A Non-Euclidean distance is based on properties of points, but not their location in a space
 - Cosine similarity, **Jaccard similarity coefficient**, ...



Jaccard Similarity

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

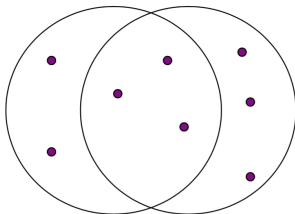
Locality-
sensitive
hashing

The Jaccard Similarity of two sets is the size of their intersection over the size of their union.

$$\text{Sim}(C_1, C_2) = \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$

The Jaccard Distance between sets is 1 minus their Jaccard similarity.

$$d(C_1, C_2) = 1 - \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$



3 in intersection

8 in union

Jaccard similarity = $3/8$

Jaccard distance = $5/8$



Outline

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

1 Finding Similar Items

2 Shingles

3 Minhashing

4 Locality-sensitive hashing



Finding Similar Documents

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Goal:

Given a large number (N in the millions or billions) of text documents, find pairs that are **near duplicates**

Applications:

- Mirror websites, or approximate mirrors
 - Don't want to show both in a search
- Similar news articles at many news sites
 - Cluster articles by **same story**

Problems:

- Many small pieces of one doc can appear out of order in another
- Too many docs to compare all pairs
- Docs are so large or so many that they cannot fit in main memory



Three Essential Steps

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- 1 **Shingling:** Convert documents, emails, etc., to sets;
- 2 **Minhashing:** Convert large sets to short signatures, while preserving similarity;
 - Depends on the distance metric;
- 3 **Locality-sensitive hashing:** Focus on pairs of signatures likely to be from similar documents.



The Big Picture

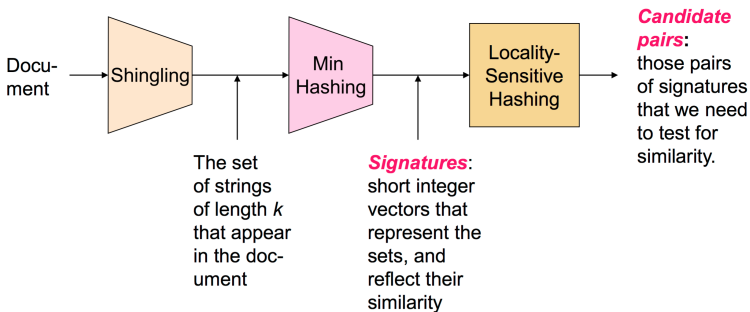
Processamento
e Recuperação
de Informação


Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Outline

Locality-sensitive hashing

- 1 Finding Similar Items
- 2 Shingles
- 3 Minhashing
- 4 Locality-sensitive hashing



Documents as High Dimensional Data

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Step 1:

Shingling: Convert documents, emails, etc., to sets

- Simple approaches...
 - Document = set of words appearing in document
 - Document = set of *important* words
- ...don't work well for this application!
 - Need to account for ordering of words
- A different way: **Shingles**



Shingles

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- A k -shingle (or k -gram) for a document is a sequence of k tokens that appears in the document
 - Tokens can be characters, words or something else, depending on application
 - Assume tokens = characters for next examples

Example: $k = 2$; $D_1 = abcab$

Set of 2-shingles: $S(D_1) = \{ab, bc, ca\}$

Option: Shingles as a bag (i.e., multi-set), counting ab twice

- Represent a doc by the set of hash values of its k -shingles



Compressing Shingles

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- To compress long shingles, we can hash them into a convenient/efficient representation (e.g. 4 bytes)
- Represent a doc by the set of hash values of its k -shingles
- **Idea:** Two documents could (rarely) appear to have shingles in common, when in fact only the hash-values were shared

Example: $k = 2$; $D_1 = abcab$

Set of 2-shingles: $S(D_1) = \{ab, bc, ca\}$

Hash the shingles: $h(D_1) = \{1, 5, 7\}$



Similarity Metric for Shingles

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Document D_1 = set of k -shingles $C_1 = S(D_1)$
- Equivalently, each document is a 0/1 vector in the space of k -shingles
 - Each unique shingle is a dimension
 - Vectors are very sparse
- A natural similarity measure is the Jaccard similarity:

$$Sim(C_1, C_2) = \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|}$$



Working Assumption

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Documents that have lots of shingles in common have similar text, even if the text appears in different order
- **Careful:** You must pick k large enough, or most documents will have most shingles
 - $k = 5$ is OK for short documents
 - $k = 10$ is better for long documents



Motivation for Minhash/LSH

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Suppose we need to find near-duplicate documents among $N = 1$ million documents
- Naively, we'd have to compute pairwise Jaccard similarites for every pair of docs
 - i.e, $\frac{N \times (N-1)}{2} \approx 5 \times 10^{11}$ comparisons
 - At 10^5 secs/day and 10^6 comparisons/sec, it would take 5 days to compute all pairwise Jaccard similarites
- For $N = 10$ million, it takes more than a year...



Outline

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- 1 Finding Similar Items
- 2 Shingles
- 3 Minhashing**
- 4 Locality-sensitive hashing



The Big Picture

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Encoding Sets as Bit Vectors

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Many similarity problems can be formalized as finding subsets that have significant intersection

- Encode sets using 0/1 (bit, Boolean) vectors
- One dimension per element in the universal set
- Interpret set intersection as bitwise AND, and set union as bitwise OR

Example: $C_1 = 10111$; $C_2 = 10011$

- Size of intersection = 3;
- Size of union = 4
- Jaccard similarity (not distance) = $3/4$
- $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 1/4$



From Sets to Boolean Matrices

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Rows = elements of the universal set
- Columns = sets
- 1 in row e and column s if and only if e is a member of s
- Column similarity is the Jaccard similarity of the sets of their rows with 1
- Typical matrix is sparse

shingles

1	0	1	0
1	1	0	1
0	1	0	1
0	0	0	1
0	0	0	1
1	1	1	0
1	0	1	0

documents



Jaccard of Columns

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Each document is a column:

Example: $C_1 = 1100011$; $C_2 = 0110010$

- Size of intersection = 2;
- size of union = 5,
- Jaccard similarity = $2/5$
- $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 3/5$

We might not really represent the data by a Boolean matrix

- Sparse matrices can be represented by the list of places with non-zero values

shingles

1	0	1	0
1	1	0	1
0	1	0	1
0	0	0	1
0	0	0	1
1	1	1	0
1	0	1	0

documents



Finding Similar Columns

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

So far:

- Documents represented as sets of shingles
- Represent sets as boolean vectors in a matrix

Next Goal: Find similar columns

- 1 Signatures of columns: small summaries of columns
- 2 Examine pairs of signatures to find similar columns
 - Essential that similarities of signatures and columns are related
- 3 **Optional:** check that columns with similar signatures are really similar

Warnings:

- Comparing all pairs may take too much time: [job for LSH](#)
- These methods can produce false negatives, and even false positives (if the optional check is not made)



Signatures of Columns

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Key idea:

Hash each column C to a small signature $h(C)$, such that:

- 1 $h(C)$ is small enough that the signature fits in RAM
- 2 $\text{sim}(C_1, C_2)$ is the same as the similarity of signatures $h(C_1)$ and $h(C_2)$

Goal:

Find a hash function $h()$ such that:

- if $\text{sim}(C_1, C_2)$ is high, then with high prob. $h(C_1) = h(C_2)$
- if $\text{sim}(C_1, C_2)$ is low, then with high prob. $h(C_1) \neq h(C_2)$

Hash docs into buckets, and expect that most pairs of near duplicate documents hash into the same bucket



Min-Hashing (1)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Goal:

Find a hash function $h()$ such that:

- if $\text{sim}(C_1, C_2)$ is high, then with high prob. $h(C_1) = h(C_2)$
- if $\text{sim}(C_1, C_2)$ is low, then with high prob. $h(C_1) \neq h(C_2)$

Clearly, the hash function depends on the similarity metric:

- Not all similarity metrics have a suitable hash function
- There is a suitable hash function for Jaccard similarity:

Min-hashing



Min-Hashing (2)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Imagine the rows of the boolean matrix permuted under random permutation π
- Define a hash function $h_{\pi}(C)$ = the number of the first (in the permuted order π) row in which column C has value 1:

$$h_{\pi}(C) = \min_{\pi}(C)$$

- Use several (e.g., 100) independent hash functions (i.e., permutations) to create a signature of a column



Min-Hashing (3)


Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Surprising Property

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Under a random permutation π ,
 $Pr[h_\pi(C_1) = h_\pi(C_2)] = sim(C_1, C_2)$

Sketch of proof:

- Let X be a set of shingles, and let $x \in X$
- Then: $Pr[\pi(y) = \min(\pi(X))] = \frac{1}{|X|}$
 - It is equally likely that any $y \in X$ is mapped to the min element
- Let x be s.t. $\pi(x) = \min(\pi(C_1 \cup C_2))$
 - Then either $\pi(x) = \min(\pi(C_1))$ if $x \in C_1$, or $\pi(x) = \min(\pi(C_2))$ if $x \in C_2$
 - So the probability that both are true is the probability $x \in C_1 \cap C_2$
- $Pr[\min(\pi(C_1)) = \min(\pi(C_2))] = \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|} = sim(C_1, C_2)$



Similarity for Signatures

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- We know $Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = sim(C_1, C_2)$
- Now generalize to multiple hash functions
- The similarity of two signatures is the fraction of the hash functions in which they agree
- **Note:** Because of the minhash property, the similarity of columns is the same as the expected similarity of their signatures



Example

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Permutation π

2	4	3
3	2	4
7	1	7
6	3	2
1	6	6
5	7	1
4	5	5

Input matrix (Shingles x Documents)

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0

Signature matrix M

2	3	7	2
2	4	1	4
7	3	7	3



Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.33	0.67	0	0



Minhash Signatures

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- Pick 100 random permutations of the rows
- Think of $\text{sig}(C)$ as a column vector
- Let $\text{sig}(C)[i] =$ according to the i -th permutation, the index of the first row that has a 1 in column C

$$\text{sig}(C)[i] = \min(\pi_i(C))$$

- **Note:** The sketch (signature) of document C is small – $\tilde{100}$ bytes!
- We achieved the goal of **compressing** long bit vectors into short signatures



Outline

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

- 1 Finding Similar Items
- 2 Shingles
- 3 Minhashing
- 4 Locality-sensitive hashing



The Big Picture

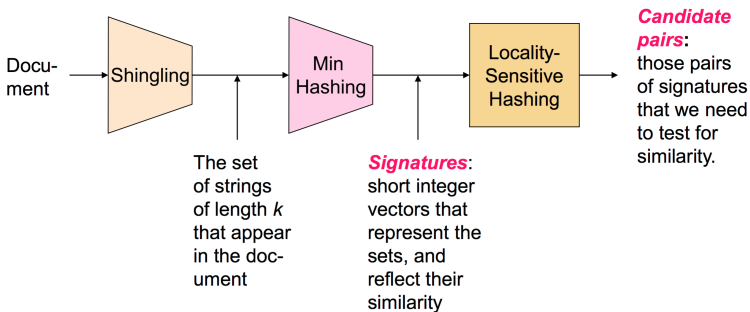
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





LSH : General Intuition

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Goal: Find documents with Jaccard similarity at least s (for some similarity threshold, e.g., $s=0.8$) LSH – Use a function $f(x,y)$ that tells whether x and y is a candidate pair, i.e. a pair of elements whose similarity must be evaluated For minhash matrices: Hash columns of signature matrix M to many buckets Each pair of documents that hashes into the same bucket is a candidate pair



Candidates from Minhash

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Pick a similarity threshold s , a fraction < 1 Columns x and y of M are a candidate pair if their signatures agree on at least fraction s of their rows:

$M(i, x) = M(i, y)$ for at least fraction s values of i

We expect documents x and y to have the same similarity as their signatures



LSH for Minhash Signatures

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Big idea: Hash columns of signature matrix M several times
Arrange that (only) similar columns are likely to hash to the
same bucket, with high probability Candidate pairs are those
that hash to the same bucket



Partition M into Bands (1)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Divide matrix M into b bands of r rows For each band, hash its portion of each column to a hash table with k buckets Make k as large as possible Candidate column pairs are those that hash to the same bucket for 1 band or more Tune b and r to catch most similar pairs, but few non-similar pairs



Partition M into Bands (2)

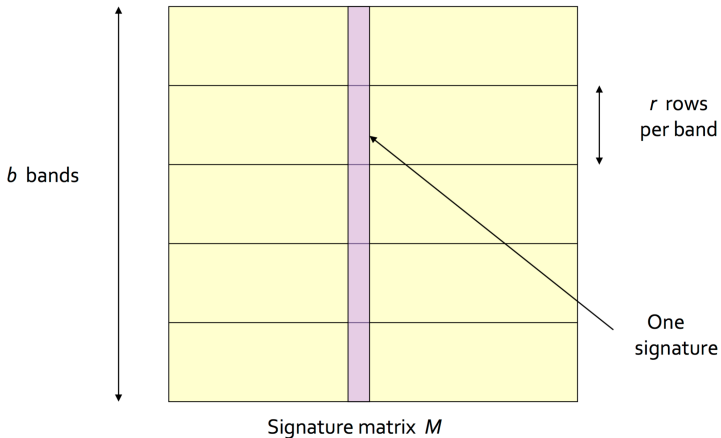
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Hashing Bands

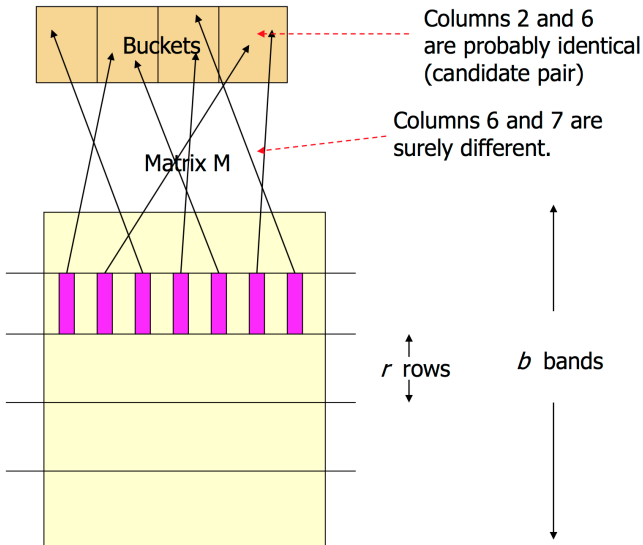
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Simplifying Assumption

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

There are enough buckets that columns are unlikely to hash to the same bucket unless they are identical in a particular band
Hereafter, we assume that **same bucket** means **identical in that band** Assumption needed only to simplify analysis, not for correctness of algorithm



Example of Bands (1)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Assume the following case: Suppose 100,000 columns of M (100k docs) Signatures of 100 integers (rows) Therefore, signatures take 40Mb Choose 20 bands of 5 integers/band
Goal: Find pairs of documents that are at least $s = 80\%$ similar



Example of Bands (2)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Assume: C1, C2 are 80% similar Since $s=80\%$ Probability C1, C2 identical in one particular band: $(0.8)^5 = 0.328$ Probability C1, C2 are not similar in all of the 20 bands: $(1 - 0.328)^{20} = 0.00035$ i.e., about 1/3000th of the 80%-similar column pairs are false negatives We would find 99.965% pairs of truly similar documents



Example of Bands (3)

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Assume: C1, C2 are 30% similar Since $s=80$ Probability C1, C2 identical in one particular band: $(0.3)^5 = 0.00243$ Probability C1, C2 identical in at least 1 of 20 bands:
 $1 - (1 - 0.00243)^{20} = 0.0474$ In other words, approximately 4.74% pairs of docs with similarity 30% end up becoming candidate pairs – false positives



LSH Involves a Tradeoff

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Pick: the number of minhashes (rows of M) the number of bands b , and the number of rows r per band to balance false positives/negatives Example: if we had only 15 bands of 5 rows, the number of false positives would go down, but the number of false negatives would go up



Analysis of LSH - What We Want

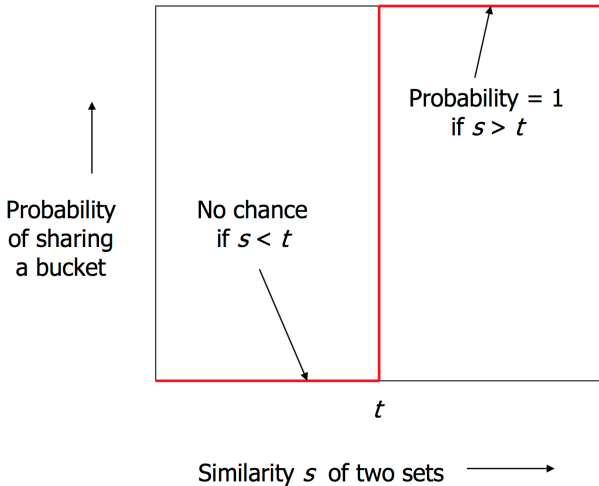
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Analysis of LSH - What One Band With One Row Gives

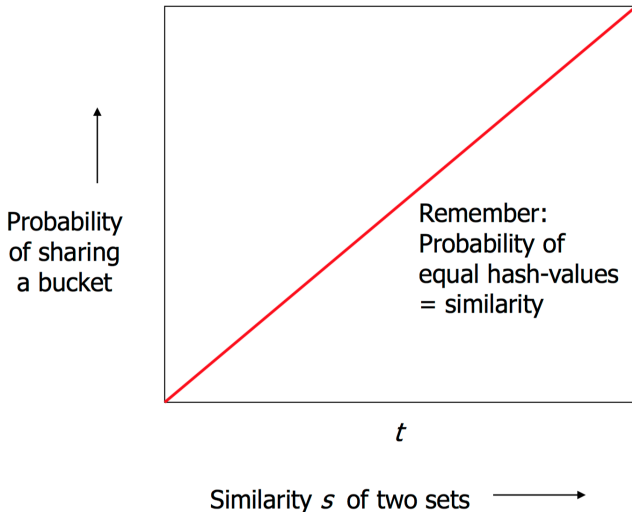
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Analysis of LSH - b bands with b rows/band

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Columns C_1 and C_2 have similarity s Pick any band (r rows)
Prob. that all rows in band equal $= s^r$ Prob. that some row in
band unequal $= 1 - s^r$ Prob. that no band identical $=$
 $(1 - s^r)^b$ Prob. that at least 1 band identical $= 1 - (1 - s^r)^b$



Analysis of LSH - What b Bands With r Rows Gives


Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





False Positives vs. False Negatives

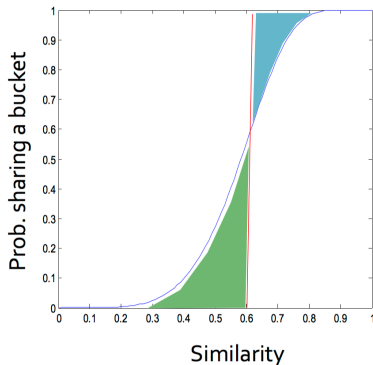
Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing



Blue area: False Negative rate

Green area: False Positive rate



LSH Summary

Processamento
e Recuperação
de Informação

Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing

Tune to get almost all pairs with similar signatures, but eliminate most pairs that do not have similar signatures Check in main memory that candidate pairs really do have similar signatures Optional: In another pass through data, check that the remaining candidate pairs really represent similar documents



The Big Picture

Processamento
e Recuperação
de Informação

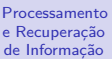
Finding
Similar Items

Shingles

Minhashing

Locality-
sensitive
hashing





Finding Similar Items

Shingles

Minhashing

Locality-sensitive hashing

Questions?